

✓ Importing libraries

```
1 from __future__ import print_function
2 import pickle
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from sklearn.metrics import classification_report
8 from sklearn import metrics
9 from sklearn import tree
10 import warnings
11 from sklearn.model_selection import train_test_split
12 from sklearn.tree import DecisionTreeClassifier
13 from sklearn.naive_bayes import GaussianNB
14 from sklearn.svm import SVC
15 from sklearn.linear_model import LogisticRegression
16 from sklearn.ensemble import RandomForestClassifier
17 from sklearn.model_selection import cross_val_score
18 warnings.filterwarnings('ignore')
```

✓ Reading Data

```
1 PATH = '/content/Crop_recommendation.csv'
2 df = pd.read_csv(PATH)
```

✓ Exploratory Data

```
1 df.head()
```

| | N | P | K | temperature | humidity | ph | rainfall | label |
|---|----|----|----|-------------|-----------|----------|------------|-------|
| 0 | 90 | 42 | 43 | 20.879744 | 82.002744 | 6.502985 | 202.935536 | rice |
| 1 | 85 | 58 | 41 | 21.770462 | 80.319644 | 7.038096 | 226.655537 | rice |
| 2 | 60 | 55 | 44 | 23.004459 | 82.320763 | 7.840207 | 263.964248 | rice |
| 3 | 74 | 35 | 40 | 26.491096 | 80.158363 | 6.980401 | 242.864034 | rice |
| 4 | 78 | 42 | 42 | 20.130175 | 81.604873 | 7.628473 | 262.717340 | rice |

```
1 df.size
```

```
17600
```

```
1 df.shape
```

```
(2200, 8)
```

```
1 df.columns
```

```
Index(['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall', 'label'], dtype='object')
```

```
1 df['label'].unique()
```

```
array(['rice', 'maize', 'chickpea', 'kidneybeans', 'pigeonpeas',
       'mothbeans', 'mungbean', 'blackgram', 'lentil', 'pomegranate',
       'banana', 'mango', 'grapes', 'watermelon', 'muskmelon', 'apple',
       'orange', 'papaya', 'coconut', 'cotton', 'jute', 'coffee'],
      dtype=object)
```

```
1 df.dtypes
```

```
N          int64
P          int64
K          int64
temperature float64
humidity    float64
ph          float64
rainfall    float64
label       object
dtype: object
```


```
1 df['label'].value_counts()
```

| | |
|-------------|-----|
| rice | 100 |
| maize | 100 |
| jute | 100 |
| cotton | 100 |
| coconut | 100 |
| papaya | 100 |
| orange | 100 |
| apple | 100 |
| muskmelon | 100 |
| watermelon | 100 |
| grapes | 100 |
| mango | 100 |
| banana | 100 |
| pomegranate | 100 |
| lentil | 100 |
| blackgram | 100 |
| mungbean | 100 |
| mothbeans | 100 |
| pigeonpeas | 100 |
| kidneybeans | 100 |
| chickpea | 100 |
| coffee | 100 |

Name: label, dtype: int64

```
1 sns.heatmap(df.corr(),annot=True)
```

<Axes: >



✓ Separating features and target label

```
1 features = df[['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall']]
2 target = df['label']
3 labels = df['label']
```

```
1 # Initializing empty lists to append all model's name and corresponding name
2 acc = []
3 model = []
```

✓ Splitting into train and test data

```
1 Xtrain, Xtest, Ytrain, Ytest = train_test_split(features, target, test_size = 0.2, random_state = 2)
```

Classification Algorithms:

▼ Decision Tree

```

1 DecisionTree = DecisionTreeClassifier(criterion="entropy",random_state=2,max_depth=5)
2 DecisionTree.fit(Xtrain,Ytrain)
3 predicted_values = DecisionTree.predict(Xtest)
4 x = metrics.accuracy_score(Ytest, predicted_values)
5 acc.append(x)
6 model.append('Decision Tree')
7 print("DecisionTrees's Accuracy is: ", x*100)
8 print(classification_report(Ytest,predicted_values))

```

```

DecisionTrees's Accuracy is: 90.0
              precision    recall  f1-score   support

   apple         1.00        1.00        1.00        13
   banana         1.00        1.00        1.00        17
 blackgram         0.59        1.00        0.74        16
  chickpea         1.00        1.00        1.00        21
   coconut         0.91        1.00        0.95        21
    coffee         1.00        1.00        1.00        22
    cotton         1.00        1.00        1.00        20
    grapes         1.00        1.00        1.00        18
     jute         0.74        0.93        0.83        28
 kidneybeans        0.00        0.00        0.00        14
    lentil         0.68        1.00        0.81        23
     maize         1.00        1.00        1.00        21
     mango         1.00        1.00        1.00        26
 mothbeans         0.00        0.00        0.00        19
  mungbean         1.00        1.00        1.00        24
 muskmelon         1.00        1.00        1.00        23
    orange         1.00        1.00        1.00        29
    papaya         1.00        0.84        0.91        19
 pigeonpeas        0.62        1.00        0.77        18
 pomegranate        1.00        1.00        1.00        17
      rice         1.00        0.62        0.77        16
 watermelon         1.00        1.00        1.00        15

 accuracy                   0.90        440
 macro avg         0.84        0.88        0.85        440
 weighted avg         0.86        0.90        0.87        440

```

```
1 # Cross validation score (Decision Tree)
2 score = cross_val_score(DecisionTree, features, target,cv=5)
```

```
1 score
```

```
array([0.93636364, 0.90909091, 0.91818182, 0.87045455, 0.93636364])
```

✓ Saving trained Decision Tree model

```
1 # Dump the trained Naive Bayes classifier with Pickle
2 DT_pkl_filename = 'DecisionTree.pkl'
3 # Open the file to save as pkl file
4 DT_Model_pkl = open(DT_pkl_filename, 'wb')
5 pickle.dump(DecisionTree, DT_Model_pkl)
6 # Close the pickle instances
7 DT_Model_pkl.close()
```

✓ Guassian Naive Bayes

```
1 NaiveBayes = GaussianNB()
2 NaiveBayes.fit(Xtrain,Ytrain)
3 predicted_values = NaiveBayes.predict(Xtest)
4 x = metrics.accuracy_score(Ytest, predicted_values)
5 acc.append(x)
6 model.append('Naive Bayes')
7 print("Naive Bayes's Accuracy is: ", x)
8 print(classification_report(Ytest,predicted_values))
```

```
Naive Bayes's Accuracy is: 0.990909090909091
      precision    recall  f1-score   support

   apple         1.00      1.00      1.00        13
  banana         1.00      1.00      1.00        17
blackgram         1.00      1.00      1.00        16
 chickpea         1.00      1.00      1.00        21
   coconut        1.00      1.00      1.00        21
    coffee         1.00      1.00      1.00        22
    cotton         1.00      1.00      1.00        20
    grapes         1.00      1.00      1.00        18
      jute         0.88      1.00      0.93        28
kidneybeans        1.00      1.00      1.00        14
    lentil         1.00      1.00      1.00        23
    maize         1.00      1.00      1.00        21
    mango         1.00      1.00      1.00        26
```

| | | | | |
|--------------|------|------|------|-----|
| mothbeans | 1.00 | 1.00 | 1.00 | 19 |
| mungbean | 1.00 | 1.00 | 1.00 | 24 |
| muskmelon | 1.00 | 1.00 | 1.00 | 23 |
| orange | 1.00 | 1.00 | 1.00 | 29 |
| papaya | 1.00 | 1.00 | 1.00 | 19 |
| pigeonpeas | 1.00 | 1.00 | 1.00 | 18 |
| pomegranate | 1.00 | 1.00 | 1.00 | 17 |
| rice | 1.00 | 0.75 | 0.86 | 16 |
| watermelon | 1.00 | 1.00 | 1.00 | 15 |
| accuracy | | | 0.99 | 440 |
| macro avg | 0.99 | 0.99 | 0.99 | 440 |
| weighted avg | 0.99 | 0.99 | 0.99 | 440 |

```

1 # Cross validation score (NaiveBayes)
2 score = cross_val_score(NaiveBayes, features, target, cv=5)
3 score

```

```
array([0.99772727, 0.99545455, 0.99545455, 0.99545455, 0.99090909])
```

▼ Saving trained Guassian Naive Bayes model

```

1 # Dump the trained Naive Bayes classifier with Pickle
2 NB_pkl_filename = 'NBClassifier.pkl'
3 # Open the file to save as pkl file
4 NB_Model_pkl = open(NB_pkl_filename, 'wb')
5 pickle.dump(NaiveBayes, NB_Model_pkl)
6 # Close the pickle instances
7 NB_Model_pkl.close()

```

▼ Support Vector Machine (SVM)

```

1
2
3 SVM = SVC(gamma='auto')
4
5 SVM.fit(Xtrain,Ytrain)
6
7 predicted_values = SVM.predict(Xtest)
8
9 x = metrics.accuracy_score(Ytest, predicted_values)
10 acc.append(x)
11 model.append('SVM')
12 print("SVM's Accuracy is: ", x)
13
14 print(classification_report(Ytest,predicted_values))

```

SVM's Accuracy is: 0.10681818181818181

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| apple | 1.00 | 0.23 | 0.38 | 13 |
| banana | 1.00 | 0.24 | 0.38 | 17 |
| blackgram | 1.00 | 0.19 | 0.32 | 16 |
| chickpea | 1.00 | 0.05 | 0.09 | 21 |
| coconut | 1.00 | 0.05 | 0.09 | 21 |
| coffee | 0.00 | 0.00 | 0.00 | 22 |
| cotton | 1.00 | 0.05 | 0.10 | 20 |
| grapes | 1.00 | 0.06 | 0.11 | 18 |
| jute | 1.00 | 0.07 | 0.13 | 28 |
| kidneybeans | 0.03 | 1.00 | 0.07 | 14 |
| lentil | 0.00 | 0.00 | 0.00 | 23 |
| maize | 0.00 | 0.00 | 0.00 | 21 |
| mango | 0.00 | 0.00 | 0.00 | 26 |
| mothbeans | 0.00 | 0.00 | 0.00 | 19 |
| mungbean | 1.00 | 0.12 | 0.22 | 24 |
| muskmelon | 1.00 | 0.30 | 0.47 | 23 |
| orange | 1.00 | 0.03 | 0.07 | 29 |
| papaya | 1.00 | 0.05 | 0.10 | 19 |
| pigeonpeas | 0.00 | 0.00 | 0.00 | 18 |
| pomegranate | 1.00 | 0.12 | 0.21 | 17 |
| rice | 0.50 | 0.06 | 0.11 | 16 |
| watermelon | 1.00 | 0.13 | 0.24 | 15 |
| accuracy | | | 0.11 | 440 |
| macro avg | 0.66 | 0.13 | 0.14 | 440 |
| weighted avg | 0.66 | 0.11 | 0.13 | 440 |

```

1 # Cross validation score (SVM)
2 score = cross_val_score(SVM,features,target,cv=5)
3 score

```



```
array([0.27727273, 0.28863636, 0.29090909, 0.275      , 0.26818182])
```

✓ Logistic Regression

```
1
2
3 LogReg = LogisticRegression(random_state=2)
4
5 LogReg.fit(Xtrain,Ytrain)
6
7 predicted_values = LogReg.predict(Xtest)
8
9 x = metrics.accuracy_score(Ytest, predicted_values)
10 acc.append(x)
11 model.append('Logistic Regression')
12 print("Logistic Regression's Accuracy is: ", x)
13
14 print(classification_report(Ytest,predicted_values))
```

```
Logistic Regression's Accuracy is: 0.9522727272727273
precision    recall  f1-score   support
```

| | | | | |
|-------------|------|------|------|----|
| apple | 1.00 | 1.00 | 1.00 | 13 |
| banana | 1.00 | 1.00 | 1.00 | 17 |
| blackgram | 0.86 | 0.75 | 0.80 | 16 |
| chickpea | 1.00 | 1.00 | 1.00 | 21 |
| coconut | 1.00 | 1.00 | 1.00 | 21 |
| coffee | 1.00 | 1.00 | 1.00 | 22 |
| cotton | 0.86 | 0.90 | 0.88 | 20 |
| grapes | 1.00 | 1.00 | 1.00 | 18 |
| jute | 0.84 | 0.93 | 0.88 | 28 |
| kidneybeans | 1.00 | 1.00 | 1.00 | 14 |
| lentil | 0.88 | 1.00 | 0.94 | 23 |
| maize | 0.90 | 0.86 | 0.88 | 21 |
| mango | 0.96 | 1.00 | 0.98 | 26 |
| mothbeans | 0.84 | 0.84 | 0.84 | 19 |
| mungbean | 1.00 | 0.96 | 0.98 | 24 |
| muskmelon | 1.00 | 1.00 | 1.00 | 23 |
| orange | 1.00 | 1.00 | 1.00 | 29 |
| papaya | 1.00 | 0.95 | 0.97 | 19 |
| pigeonpeas | 1.00 | 1.00 | 1.00 | 18 |
| pomegranate | 1.00 | 1.00 | 1.00 | 17 |
| rice | 0.85 | 0.69 | 0.76 | 16 |
| watermelon | 1.00 | 1.00 | 1.00 | 15 |

```
accuracy                0.95    440
macro avg               0.95    0.95    0.95    440
```

```
weighted avg      0.95      0.95      0.95      440
```

```
1 # Cross validation score (Logistic Regression)
2 score = cross_val_score(LogReg, features, target, cv=5)
3 score
```

```
array([0.95      , 0.96590909, 0.94772727, 0.96590909, 0.94318182])
```

▼ Saving trained Logistic Regression model

```
1 # Dump the trained Naive Bayes classifier with Pickle
2 LR_pkl_filename = 'LogisticRegression.pkl'
3 # Open the file to save as pkl file
4 LR_Model_pkl = open(DT_pkl_filename, 'wb')
5 pickle.dump(LogReg, LR_Model_pkl)
6 # Close the pickle instances
7 LR_Model_pkl.close()
```

▼ Random Forest

```
1
2 RF = RandomForestClassifier(n_estimators=20, random_state=0)
3 RF.fit(Xtrain, Ytrain)
4
5 predicted_values = RF.predict(Xtest)
6
7 x = metrics.accuracy_score(Ytest, predicted_values)
8 acc.append(x)
9 model.append('RF')
10 print("RF's Accuracy is: ", x)
11
12 print(classification_report(Ytest, predicted_values))
```

```
RF's Accuracy is: 0.990909090909091
```

| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| apple | 1.00 | 1.00 | 1.00 | 13 |
| banana | 1.00 | 1.00 | 1.00 | 17 |
| blackgram | 0.94 | 1.00 | 0.97 | 16 |
| chickpea | 1.00 | 1.00 | 1.00 | 21 |
| coconut | 1.00 | 1.00 | 1.00 | 21 |
| coffee | 1.00 | 1.00 | 1.00 | 22 |

| | | | | |
|--------------|------|------|------|-----|
| cotton | 1.00 | 1.00 | 1.00 | 20 |
| grapes | 1.00 | 1.00 | 1.00 | 18 |
| jute | 0.90 | 1.00 | 0.95 | 28 |
| kidneybeans | 1.00 | 1.00 | 1.00 | 14 |
| lentil | 1.00 | 1.00 | 1.00 | 23 |
| maize | 1.00 | 1.00 | 1.00 | 21 |
| mango | 1.00 | 1.00 | 1.00 | 26 |
| mothbeans | 1.00 | 0.95 | 0.97 | 19 |
| mungbean | 1.00 | 1.00 | 1.00 | 24 |
| muskmelon | 1.00 | 1.00 | 1.00 | 23 |
| orange | 1.00 | 1.00 | 1.00 | 29 |
| papaya | 1.00 | 1.00 | 1.00 | 19 |
| pigeonpeas | 1.00 | 1.00 | 1.00 | 18 |
| pomegranate | 1.00 | 1.00 | 1.00 | 17 |
| rice | 1.00 | 0.81 | 0.90 | 16 |
| watermelon | 1.00 | 1.00 | 1.00 | 15 |
| accuracy | | | 0.99 | 440 |
| macro avg | 0.99 | 0.99 | 0.99 | 440 |
| weighted avg | 0.99 | 0.99 | 0.99 | 440 |

```

1 # Cross validation score (Random Forest)
2 score = cross_val_score(RF,features,target,cv=5)
3 score

```

```
array([0.99772727, 0.99545455, 0.99772727, 0.99318182, 0.98863636])
```

▼ Saving trained Random Forest model

```


1
2 # Dump the trained Naive Bayes classifier with Pickle
3 RF_pkl_filename = 'RandomForest.pkl'
4 # Open the file to save as pkl file
5 RF_Model_pkl = open(RF_pkl_filename, 'wb')
6 pickle.dump(RF, RF_Model_pkl)
7 # Close the pickle instances
8 RF_Model_pkl.close()

```

▼ Accuracy Comparison

```
1 plt.figure(figsize=[10,5],dpi = 100)
2 plt.title('Accuracy Comparison')
3 plt.xlabel('Accuracy')
4 plt.ylabel('Algorithm')
5 sns.barplot(x = acc,y = model,palette='dark')
```

<Axes: title={'center': 'Accuracy Comparison'}, xlabel='Accuracy', ylabel='Algorithm'>



```
1 accuracy_models = dict(zip(model, acc))
2 for k, v in accuracy_models.items():
3     print (k, '-->', v)
```

```
Decision Tree --> 0.9
Naive Bayes --> 0.990909090909091
SVM --> 0.106818181818181
Logistic Regression --> 0.952272727272727
RF --> 0.990909090909091
```

✓ Making a prediction

```
1 data = np.array([[83, 45, 60, 28, 70.3, 7.0, 150.9]])
2 prediction = RF.predict(data)
3 print(prediction)
```

```
1 data = np.array([[83, 45, 60, 28, 70.3, 7.0, 150.9]])
2 prediction = LogReg.predict(data)
3 print(prediction)
```

```
['jute']
```

```
1 data = np.array([[83, 45, 60, 28, 70.3, 7.0, 150.9]])
2 prediction = SVM.predict(data)
3 print(prediction)
```

```
['kidneybeans']
```

```
1 data = np.array([[83, 45, 60, 28, 70.3, 7.0, 150.9]])
```